

# **Consensus Algorithms - III**

(Practical Byzantine fault tolerance Algorithm  
&  
Three phase commit Protocol)

# Distributed Consensus - Properties

- **Termination:** Every correct individual decides some value at the end of the consensus protocol
- **Validity:** If all the individuals proposes the same value, then all correct individuals decide on that value
- **Integrity:** Every correct individual decides at most one value, and the decided value must be proposed by some individuals
- **Agreement:** Every correct individual must agree on the same value

# Synchronous vs Asynchronous Systems

- **Synchronous Message Passing System:** The message must be received within a predefined time interval
  - Strong guarantee on message transmission delay
- **Asynchronous Message Passing System:** There is no upper bound on the message transmission delay or the message reception time
  - No timing constraint, message can be delayed for arbitrary period of times

# Asynchronous Consensus

- **FLP85 (Impossibility Result):** In a **purely asynchronous distributed system**, the consensus problem is **impossible (with a deterministic solution)** to solve if in the presence of a **single crash failure**.
  - Results by Fischer, Lynch and Patterson (most influential paper awarded in ACM PODC 2001)
  - Randomized algorithms may exist

# Correctness of Distributed Consensus

- **Safety:**

Correct individuals must not agree on incorrect value

- Nothing bad happened

- **Liveliness (or Liveness):**

Every correct value must be accepted eventually

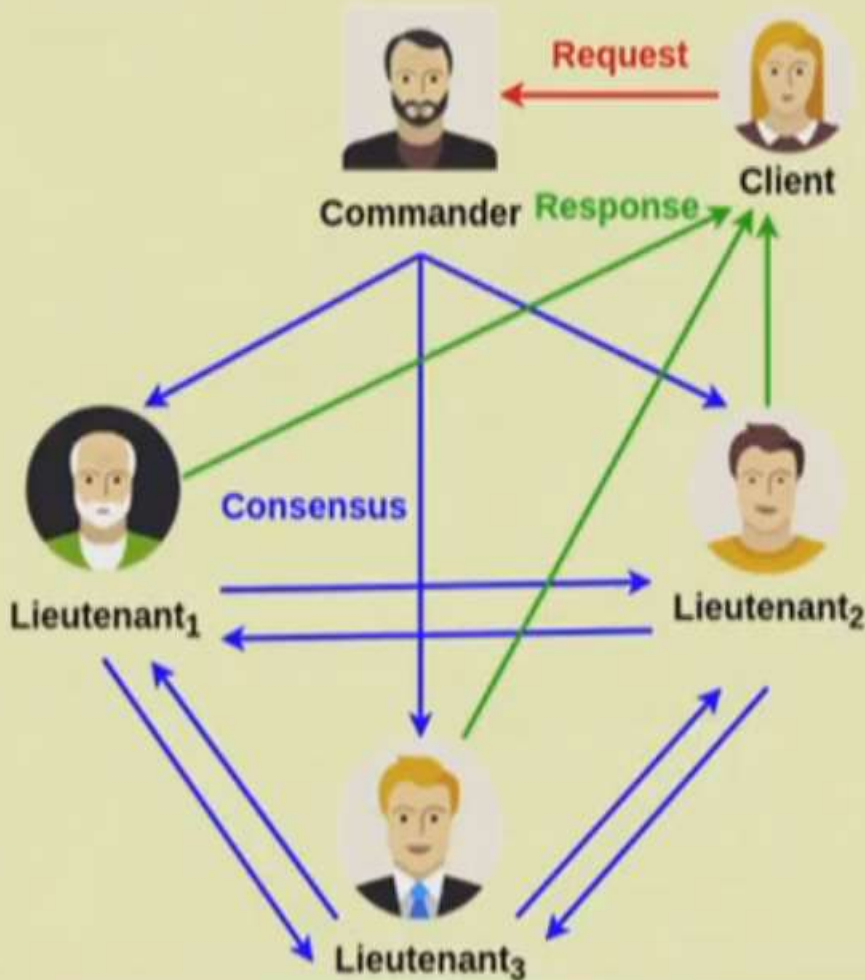
- Something good eventually happens

# Practical Byzantine Fault Tolerant

- Why **Practical**?
  - Ensures safety over an asynchronous network (not liveness!)
  - Byzantine Failure
  - Low overhead
- **Real Applications**
  - Tendermint
  - IBM's Openchain
  - ErisDB
  - Hyperledger



# Practical Byzantine Fault Tolerant Model



- Asynchronous distributed system
  - delay, out of order message
- Byzantine failure handling
  - arbitrary node behavior
- Privacy
  - tamper-proof message, authentication

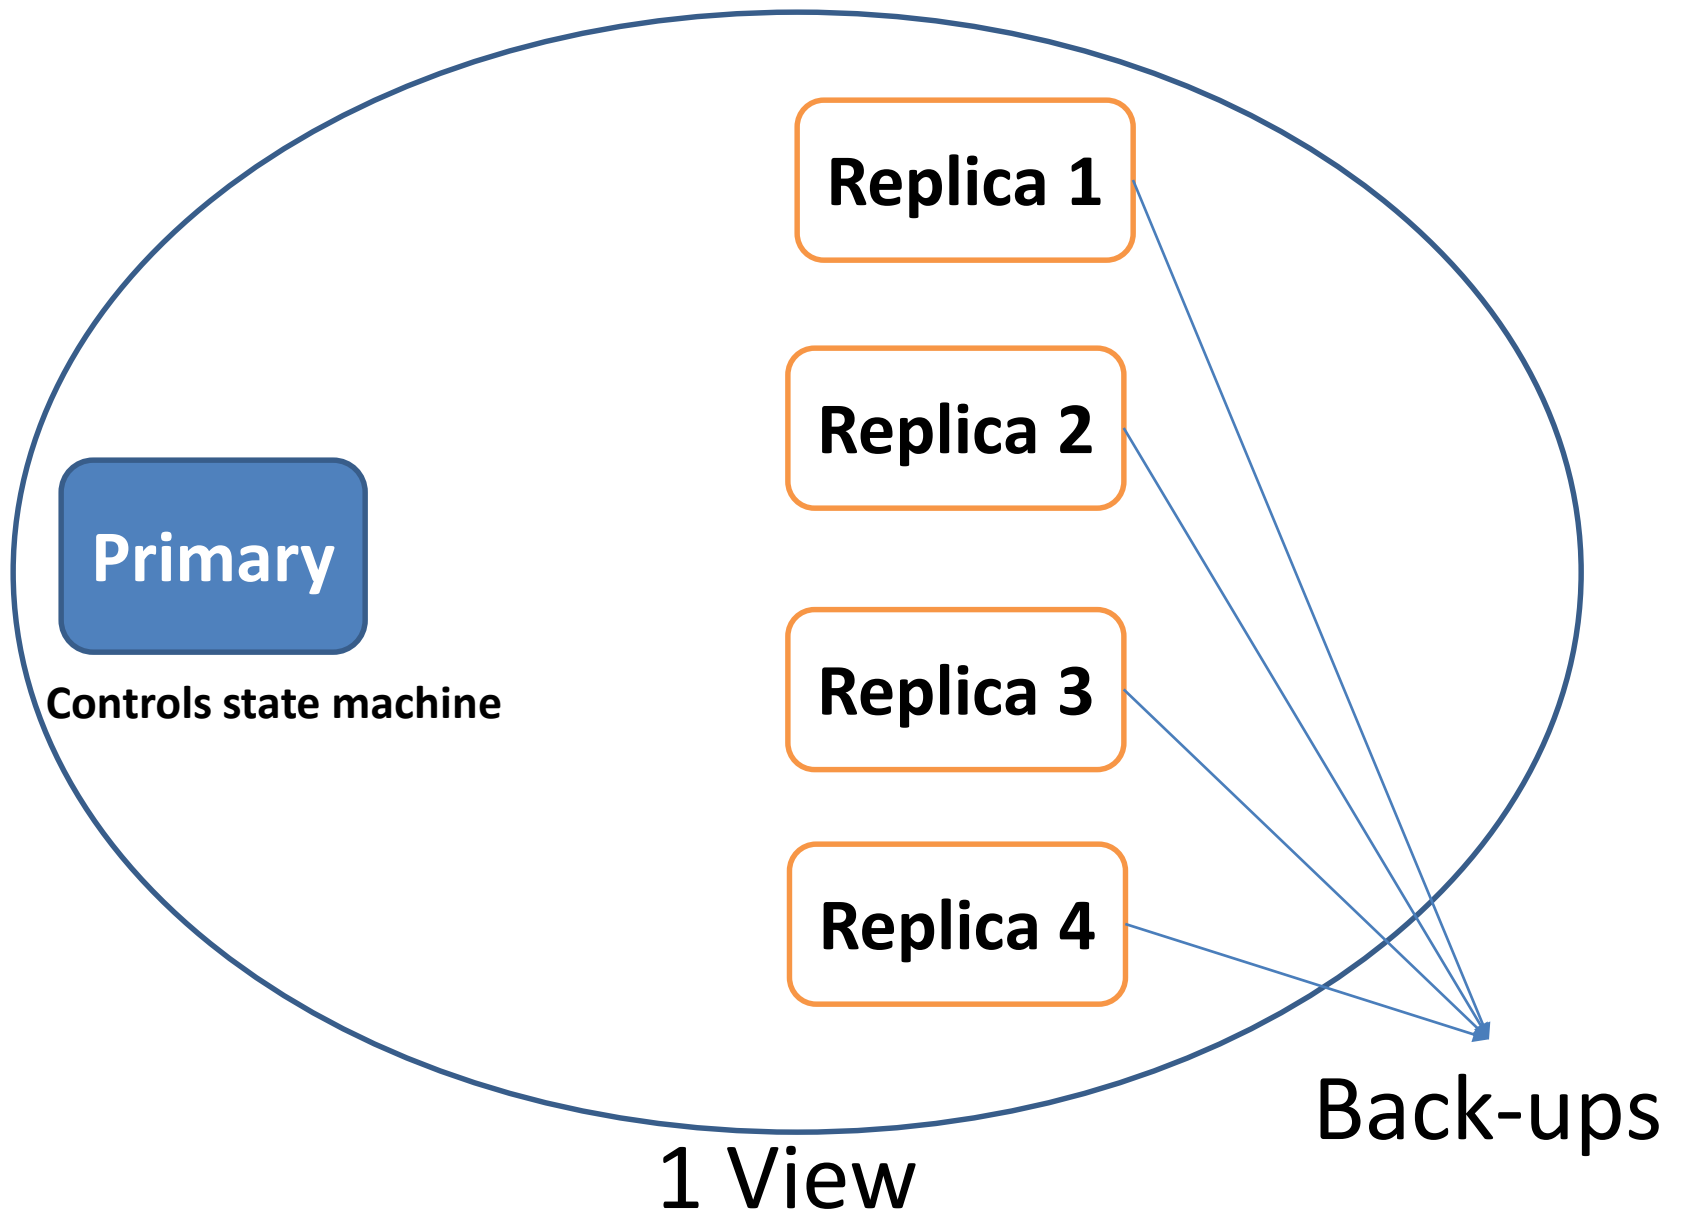
Hashing

Digital  
Signature

# Practical Byzantine Fault Tolerant Model

- A state machine is replicated across different nodes
- $3f + 1$  replicas are there where  $f$  is the number of faulty replicas
- The replicas move through a successions of configurations, known as *views*  
**States of commanders and lieutenants**
- One replica in a *view* is *primary* and others are *backups*






\*\* In another **View**, primary may become replica and one of the replica may become primary

# Practical Byzantine Fault Tolerant Model

- Views are changed when a *primary* is detected as faulty
- Every view is identified by a unique integer number  $v$
- Only the messages from the current views are accepted

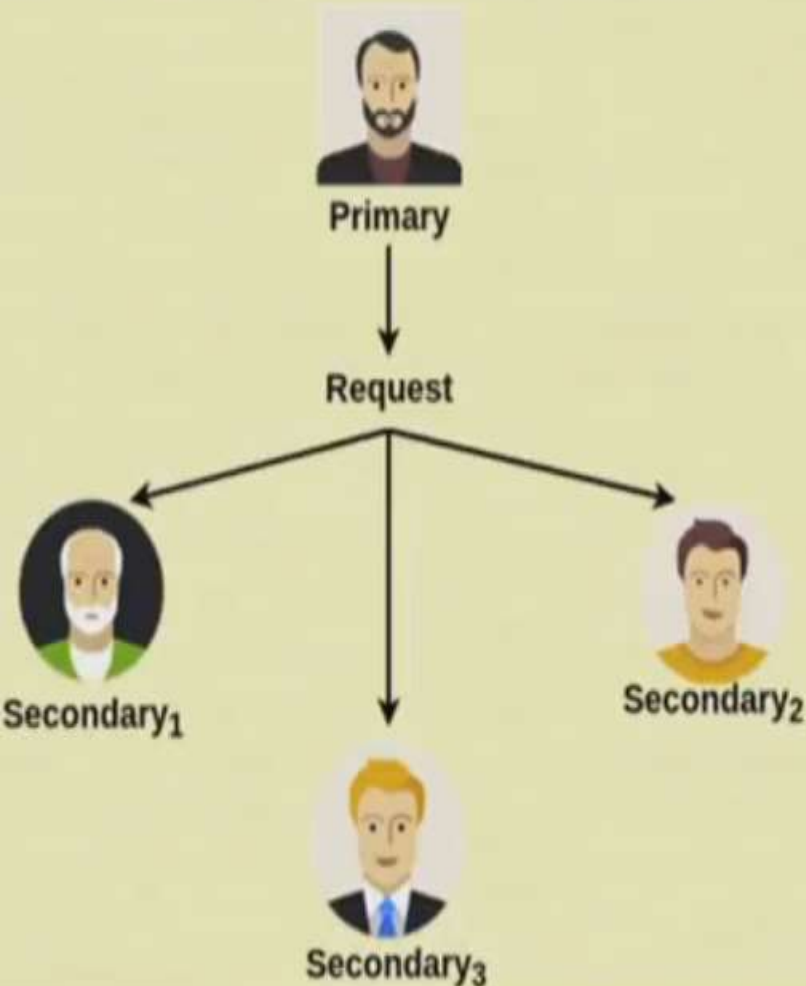
# Practical Byzantine Fault Tolerant Algorithm



- A client sends a request to invoke a service operation to the primary

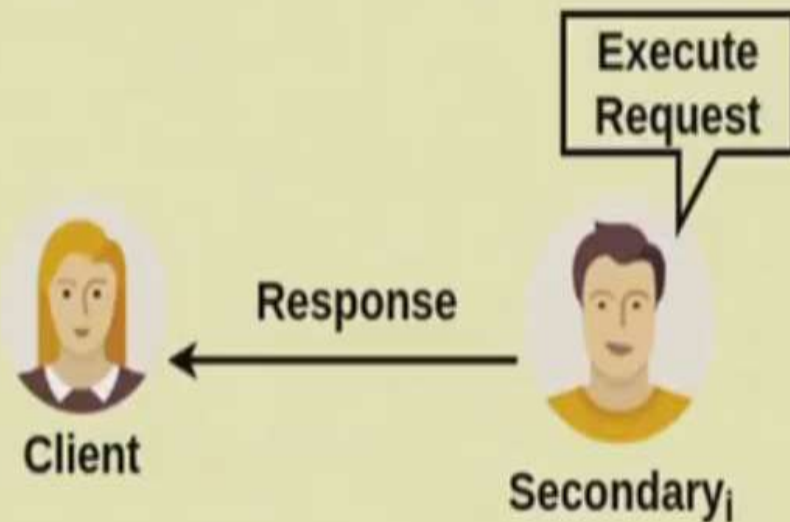
(Transaction request)

# Practical Byzantine Fault Tolerant Algorithm



- The primary multicasts the request to the backups

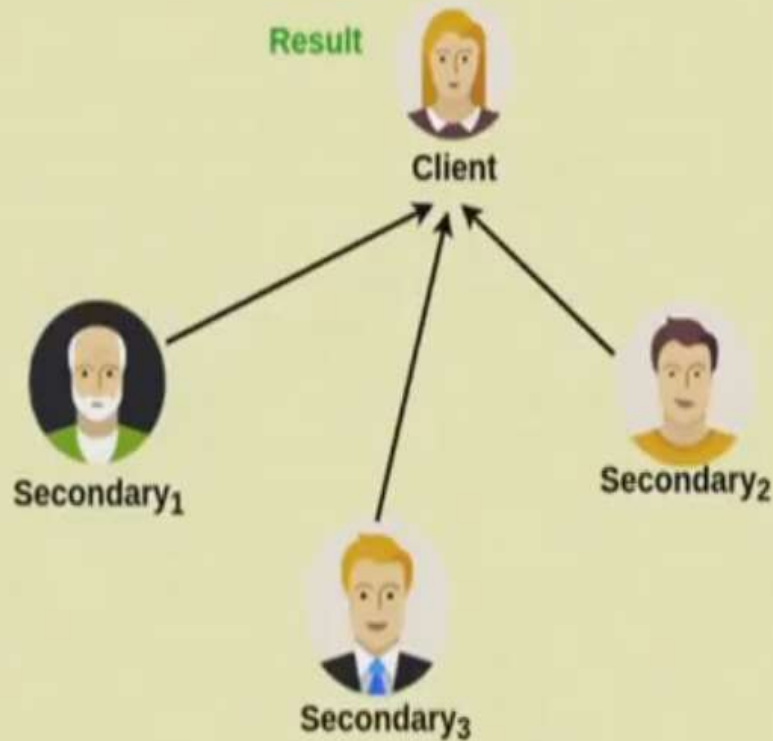
# Practical Byzantine Fault Tolerant Algorithm



- Backups execute the request and send a reply to the client



# Practical Byzantine Fault Tolerant Algorithm



- The client waits for  $f + 1$  replies from different backups with the same result
  - $f$  is the maximum number of faulty replicas that can be tolerated


**Total:  $3f + 1$**

**Faulty:  $f$**

**Non-Faulty:  $3f + 1 - f \Rightarrow 2f + 1$**

**Majority:  $f + 1$**


PBFT works in three phases: Pre-prepare, Prepare & Commit



## Three Phase Commit Protocol - Pre-Prepare

- **Pre-prepare:** Primary assigns a sequence number  $n$  to the request and multicast a message  $\langle \langle PRE - PREPARE, v, n, d \rangle_{\sigma_p}, m \rangle$  to all the backups
  - $v$  is the current view number
  - $n$  is the message sequence number
  - $d$  is the message digest
  - $\sigma_p$  is the private key of primary - works as a digital signature
  - $m$  is the message to transmit

# Three Phase Protocol




- Pre-prepare:
  - Acknowledge the request by a unique sequence number

## Three Phase Commit Protocol - Pre-Prepare

- Pre-prepare messages are used as a proof that request was assigned sequence number  $n$  is the view  $v$
- A backup accepts a pre-prepare message if
  - The signature is correct and  $d$  is the digest for  $m$
  - The backup is in view  $v$
  - It has not received a different PRE-PREPARE message with sequence  $n$  and view  $v$  with a different digest
  - The sequence number is within a threshold



# Three Phase Protocol



- Prepare:
  - Replicas agree on the assigned sequence number

## Three Phase Commit Protocol - Prepare

- If the backup accepts the PRE-PREPARE message, it enters prepare phase by multicasting a message  $\langle PREPARE, v, n, d, i \rangle_{\sigma_i}$  to all other replicas
- A replica (both primary and backups) accepts prepare messages if
  - Signatures are correct
  - View number equals to the current view
  - Sequence number is within a threshold

# Three Phase Commit Protocol

- Pre-prepare and prepare ensure that non-faulty replicas guarantee on a total order for the requests within a view
- Commit a message if
  - $2f$  prepares from different backups matches with the corresponding pre-prepare
  - You have total  $2f + 1$  votes (one from primary that you already have!) from the non-faulty replicas

## Three Phase Commit Protocol

**Why do you require  $3f + 1$  replicas to ensure safety in an asynchronous system when there are  $f$  faulty nodes?**

- If you have  $2f + 1$  replicas, you need all the votes to decide the majority - boils down to a synchronous system
- You may not receive votes from certain replicas due to delay, in case of an asynchronous system
- $f + 1$  votes do not ensure majority, may be you have received  $f$  votes from Byzantine nodes, and just one vote from a non-faulty node (note Byzantine nodes can vote for or against - You do not know that a priori!)



# Three Phase Commit Protocol


- **Why do you require  $3f + 1$  replicas to ensure safety in an asynchronous system when there are  $f$  faulty nodes?**
  - If you do not receive a vote
    - The node is faulty and not forwarded a vote at all
    - The node is non-faulty, forwarded a vote, but the vote got delayed
  - Majority can be decided once  $2f + 1$  votes have arrived - even if  $f$  are faulty, you know  $f + 1$  are from correct nodes, do not care about the remaining  $f$  votes



## Three Phase Commit Protocol - Commit

- Multicast  $\langle COMMIT, v, n, d, i \rangle_{\sigma_i}$  message to all the replicas including primary
- Commit a message when a replica
  - Has sent a commit message itself
  - Has received  $2f + 1$  commits (including its own)

# Three Phase Protocol

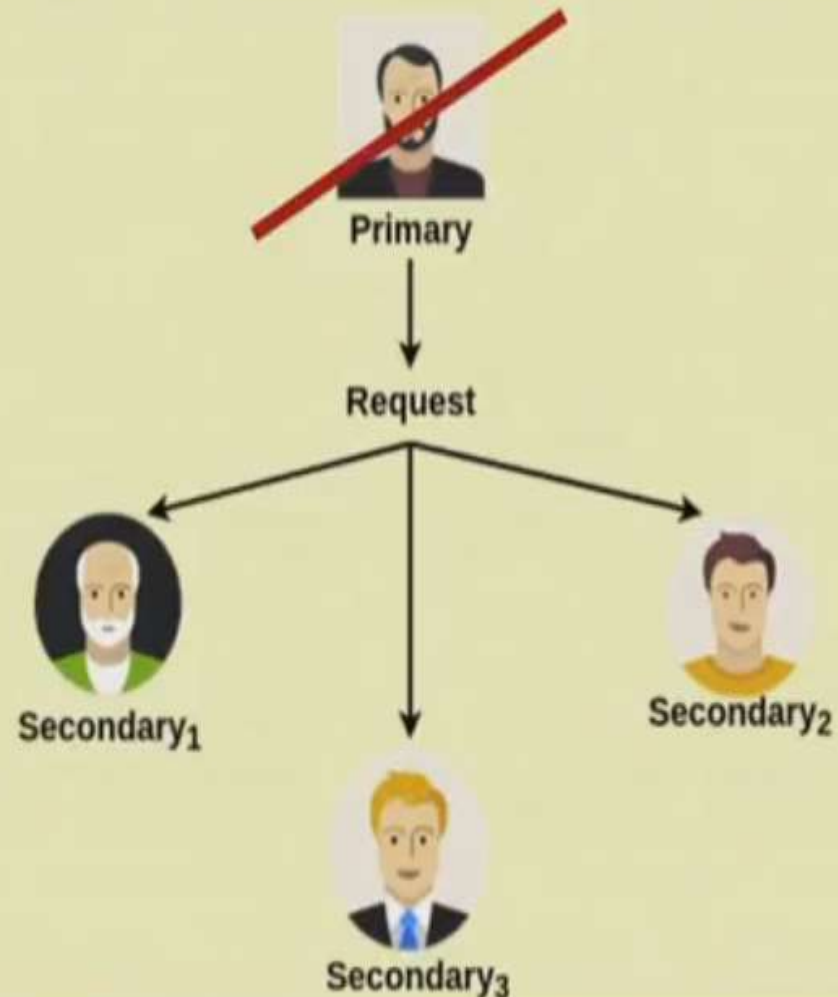


- Commit:
  - Establish consensus throughout the views

# View Change

What if the **primary** is **faulty**??

- non-faulty replicas detect the fault
- replicas together start view change operation



# Primary Faulty conditions

- **Problem (Case 1)**
- Sequence number 1: INSERT (APPLE) INTO FRUIT
- Sequence number 4: INSERT (PEAR) INTO FRUIT
- Sequence number 5: SELECT \* FROM FRUIT The
- replica will be stuck waiting for request with sequence number 2...
- **Problem (Case 2)**
  - Client sends request to primary
  - Primary doesn't forward the request to the replicas...

# View Changes

- View-change protocol provides **liveness**
  - Allow the system to make progress when primary fails
- If the primary fails, backups will not receive any message (such as PRE\_PREPARE or COMMIT) from the primary
- View changes are triggered by timeouts
  - Prevent backups from waiting indefinitely for requests to execute



# Correctness of Distributed Consensus

- **Safety:**

Correct individuals must not agree on incorrect value

- Nothing bad happened

- **Liveliness (or Liveness):**

Every correct value must be accepted eventually

- Something good eventually happens

## View Change

- Backup starts a timer when it receives a request, and the timer is not already running
  - The timer is stopped when the request is executed
  - Restarts when some new request comes
- If the timer expires at view  $v$ 
  - Backup starts a view change to move the system to view  $v + 1$

## View Change

- On timer expiry, a backup stops accepting messages except
  - Checkpoint
  - View-change
  - New-View

## View Change

Multicasts a  $\langle VIEW\_CHANGE, v + 1, n, \mathcal{C}, \mathcal{P}, i \rangle_{\sigma_i}$  message to all replicas

- $n$  is the sequence number of the last stable checkpoint  $s$  known to  $i$
- $\mathcal{C}$  is a set of  $2f + 1$  valid checkpoint messages proving the correctness of  $s$
- $\mathcal{P}$  is a set containing a set  $\mathcal{P}_m$  for each request  $m$  that prepared at  $i$  with a sequence number higher than  $n$ 
  - Each set  $\mathcal{P}_m$  contains a valid pre-prepare message and  $2f$  matching

# Correctness

- **Safety:** The algorithm provides safety if all non-faulty replicas agree on the sequence numbers of requests that commit locally





# Correctness

**Liveness:** To provide liveness, replicas must move to a new view if they are unable to execute a request

- A replica waits for  $2f + 1$  view change messages and then starts a timer to initiate a new view (*avoid starting a view change too soon*)
- If a replica receives a set of  $f + 1$  valid view change messages for views greater than its current view, it sends view change message (*prevents starting the next view change too late*)
- Faulty replicas are unable to impede progress by forcing frequent view change

## Consensus in Permissioned Model

- PBFT has well adopted in consensus for permissioned blockchain environments
  - Hyperledger
  - Tendermint Core
- Several scalability issues are still there, we'll discuss those in details in the later part of the course!

**THANK YOU**