

Research Prospective

Consensus Protocols

- Permissionless Blockchain
 - PoW
 - PoS
 - PoET
 - PoB
 - .
 - .
 - .
- Permissioned Blockchain
 - RAFT
 - PAXOS
 - BFT
 - PBFT
 - RBFT (Redundant Byzantine Fault Tolerance)

PoW Vs PBFT

- PoW

- Open environment
- Works over a large number of nodes
- Scalable in terms of number of nodes (+ve)
- Transaction throughput is low:
 - Very less no. of transactions per block (-ve)
 - No. of transactions per unit time

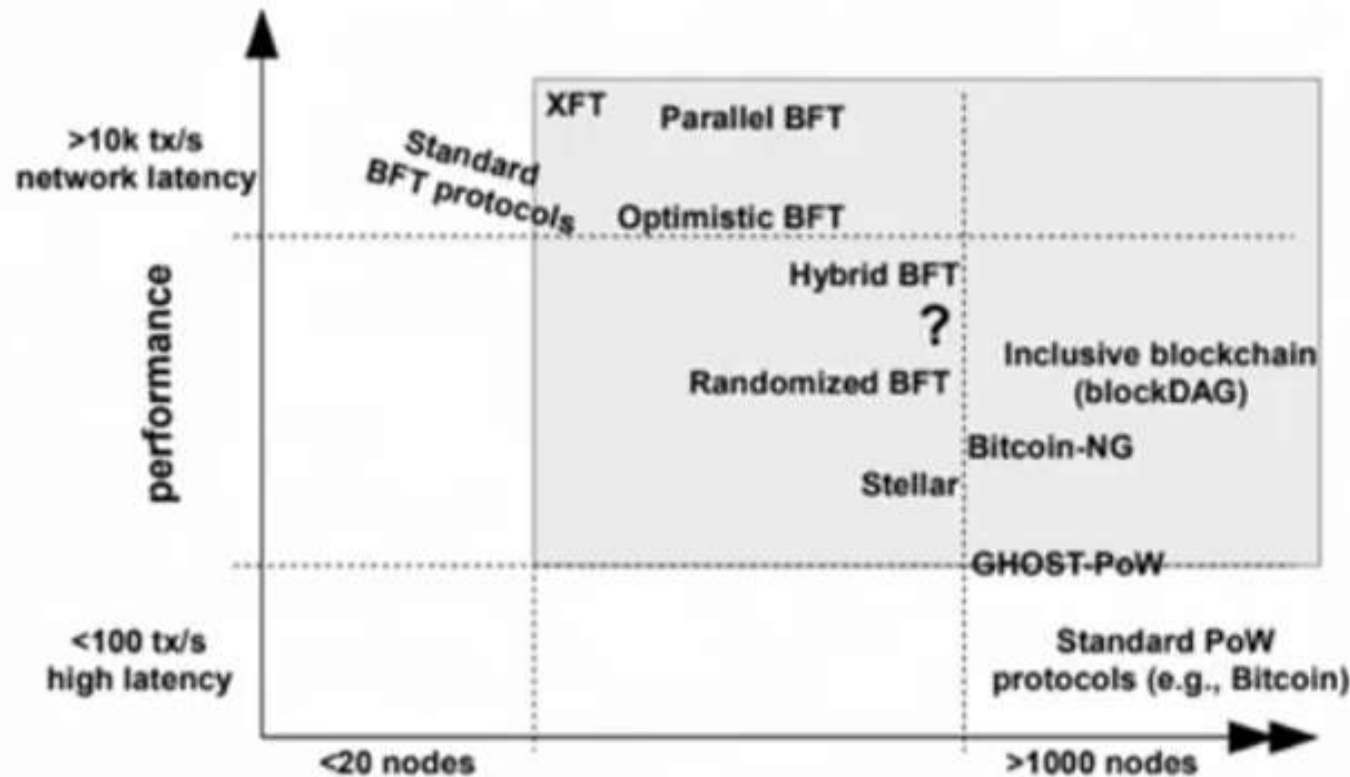
- PBFT

- Closed
- Not scalable in terms of number of nodes
- High transaction throughput (+ve)
- High message complexity (-ve)

PoW Scalability

- Two magic numbers in PoW
 - **Block frequency** - 10 minutes
 - **Block size** - 1 MB **In 2008, but later it increased upto 8 MB**
 - Transaction throughput - 7 transactions per second (with 200-250 bytes transactions)
-
- VISA or Master card transactions: About 40-45 million per second

Performance vs Scalability for PoW and BFT



Vukolić, Marko. "The quest for scalable blockchain fabric: Proof-of-work vs. BFT replication." *International Workshop on Open Problems in Network Security*. Springer, Cham, 2015.

PoW vs PBFT - Consensus Finality

- *If a correct node p appends block b to its copy of blockchain before appending block b' , then no correct node q appends block b' before b to its copy of the blockchain (Vukolic, 2015)*
- PoW is a randomized protocol - does not ensure consensus finality
 - Remember the forks in Bitcoin blockchain
- BFT protocols ensure total ordering of transactions - ensures consensus finality

	PoW consensus	BFT consensus
Node identity management	open, entirely decentralized	permissioned, nodes need to know IDs of all other nodes
Consensus finality	no	yes
Scalability (no. of nodes)	excellent (thousands of nodes)	limited, not well explored (tested only up to $n \leq 20$ nodes)
Scalability (no. of clients)	excellent (thousands of clients)	excellent (thousands of clients)
Performance (throughput)	limited (due to possible of chain forks)	excellent (tens of thousands tx/sec)
Performance (latency)	high latency (due to multi-block confirmations)	excellent (matches network latency)
Power consumption	very poor (PoW wastes energy)	good
Tolerated power of an adversary	$\leq 25\%$ computing power	$\leq 33\%$ voting power
Network synchrony assumptions	physical clock timestamps (e.g., for block validity)	none for consensus safety (synchrony needed for liveness)
Correctness proofs	no	yes

Bitcoin-NG



Eyal, I., Gencer, A. E., Sirer, E. G., & Van Renesse, R.
(2016, March). **Bitcoin-NG: A Scalable Blockchain
Protocol.** In *NSDI 2016*

Bitcoin-NG

Issues with Nakamoto Consensus (PoW)


- **Transaction scalability**
 - Block frequency of 10 minutes and block size of 1 MB during mining reduces the transactions supported per second
- **Issues with Forks**
 - Prevents consensus finality
 - Makes the system unfair - a miner with poor connectivity has always in a disadvantageous position

Bitcoin-NG: A Scalable PoW Protocol

- Bitcoin - think of the winning miner as the **leader** - the leader serializes the transactions and include a new block in the blockchain
- Decouple Bitcoin's blockchain operations into two planes
 - **Leader election:** Use PoW to randomly select a leader (an infrequent operation)
 - **Transaction Serialization:** The leader serializes the transaction until a new leader is elected

Bitcoin vs Bitcoin-NG

Bitcoin



Bitcoin-NG

Key Block

Micro Blocks



Key blocks generated by PoW

Contains key of potential miner

Bitcoin-NG: Key Blocks

- Key blocks are used to choose a leader (similar to Bitcoin)
- A key block contains
 - The reference to the previous block
 - The current Unix time
 - A coinbase transaction to pay of the reward
 - A target hash value
 - A nonce field

Bitcoin-NG: Key Blocks

- For a key block to be valid, the cryptographic hash of its header must be smaller than the target value.
- The key block also contains a public key (so the name, key block)
 - Used in subsequent microblocks



Bitcoin-NG: Key Blocks

- Key blocks are generated based on regular Bitcoin mining procedure
 - Find out the nonce such that the block hash is less than the target value
- Key blocks are generated infrequently - the intervals between two key blocks is exponentially distributed



Bitcoin-NG: Microblocks

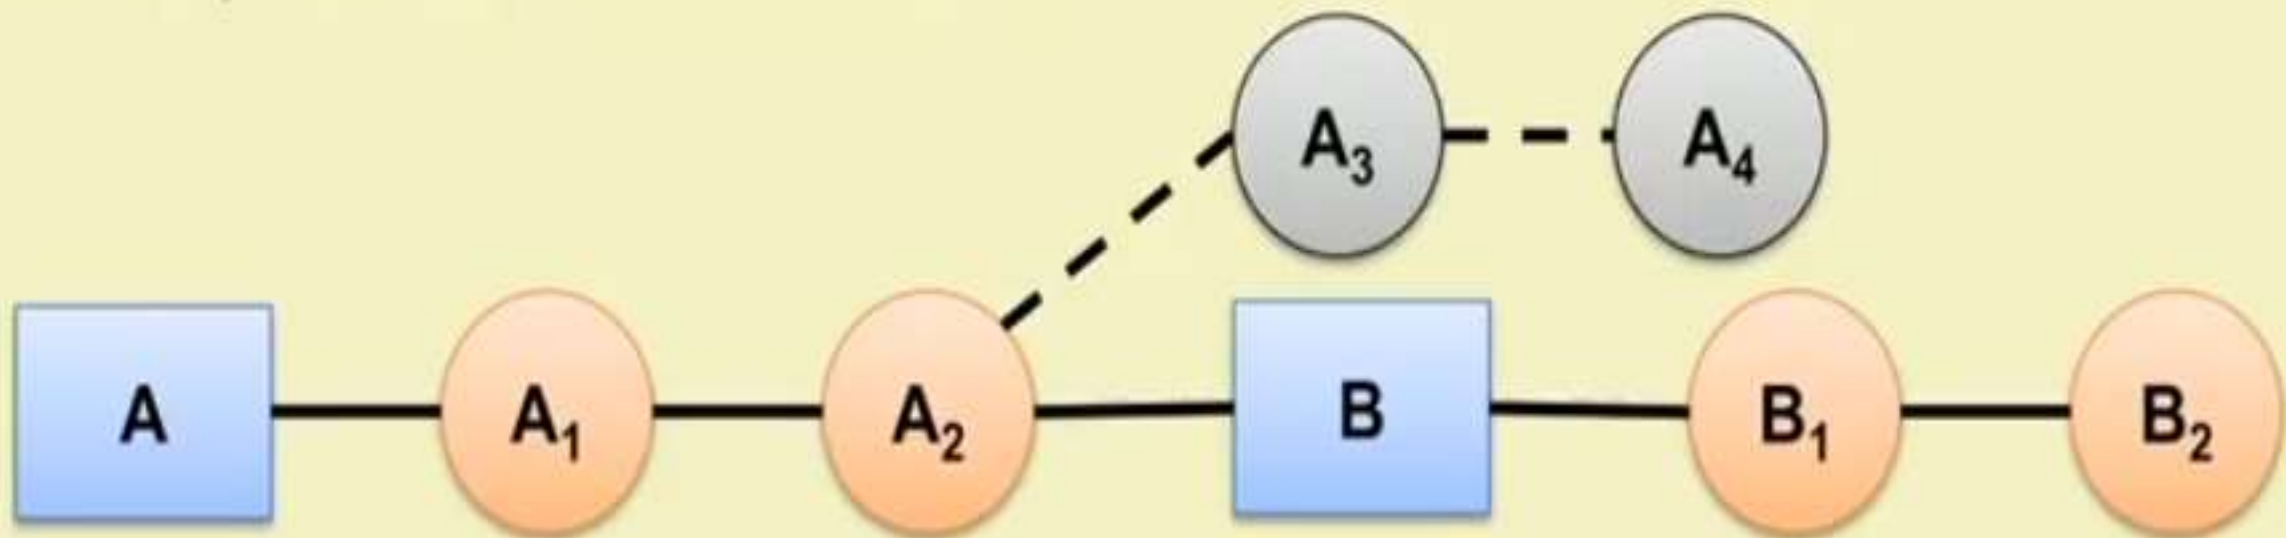
- Once a node generates a key block, it becomes the **leader**
- As a leader, the node is allowed to generate microblocks
 - Microblocks are generated at a set rate smaller than a predefined maximum
 - The rate is much higher than the key block generation rate

Bitcoin-NG: Microblocks

- A microblock contains
 - Ledger entries
 - Header
 - Reference to the previous block
 - The current Unix time
 - A cryptographic hash of the ledger entries (Markle root)
 - A cryptographic signature of the header
- The signature uses private key corresponding to the key block public key


Bitcoin-NG: Confirmation Time

- When a miner generates a key block, he may not have heard of all microblocks generated by the previous leader
 - Common if microblock generation is frequent
 - May result in microblock fork




Bitcoin-NG: Confirmation Time

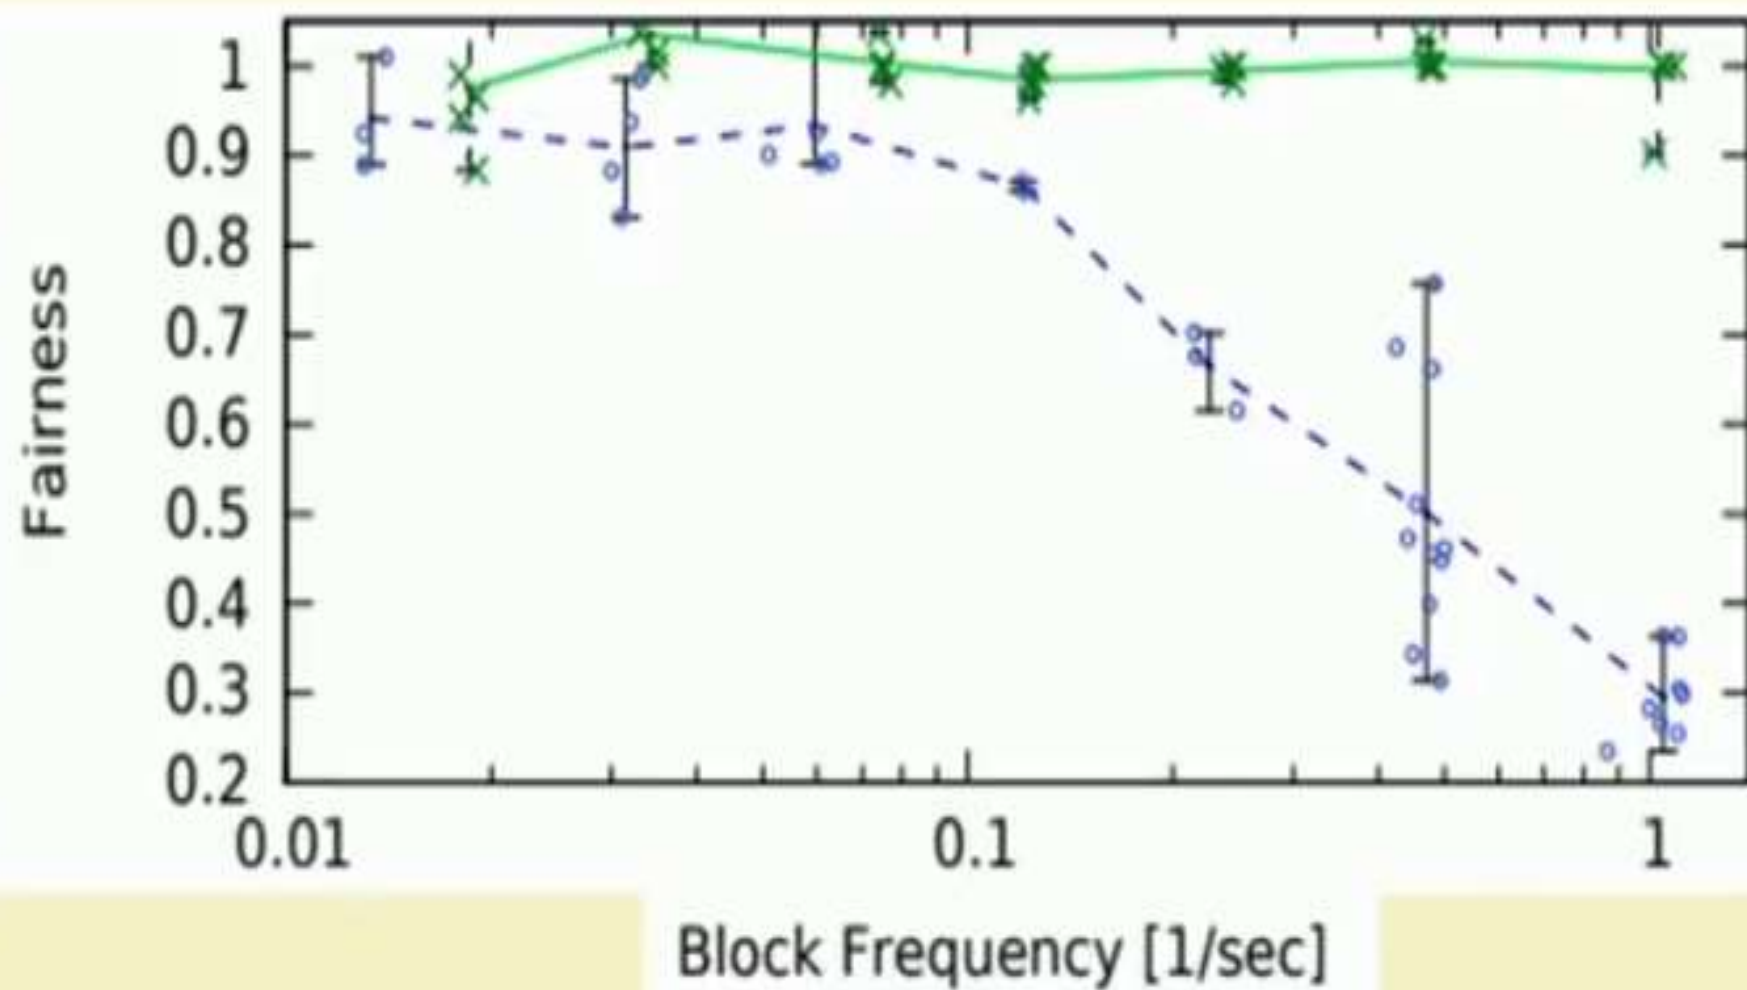
- A node may hear a forked microblock (A_3) but not the new key block (B)
 - This can be prevented by ensuring the reception of the key block
 - When a node sees a microblock, it waits for propagation time of the network, to make sure it is not pruned by a new key block



Bitcoin vs Bitcoin-NG



Bitcoin vs Bitcoin-NG



Requirements for Blockchain Consensus


- **Byzantine fault tolerant** – the system should work even in the presence of malicious users while operating across multiple administrative domains
- Should provide **strong consistency guarantee** across replicas
- Should **scale well to increasing workloads** in terms of transactions processed per unit time
- Should **scale well to increasing network size**

Collective Signing (CoSi)

- Method to protect “authorities and their clients” from undetected misuse or exploits
- A **scalable witness cosigning protocol** ensuring that every authoritative statement is validated and publicly logged by a diverse group of witnesses before any client accepts it
- A statement S collectively signed by W witnesses assures clients that S has been seen, and not immediately found erroneous, by those W observers.

CoSi Architecture

Authoritative statements: e.g. log records



- The leader organizes the witnesses in a tree structure – a scalable way of aggregating signatures coming from the children
- Three rounds of PBFT (pre-prepare, prepare and commit) can be simulated using two rounds of CoSi

CoSi Architecture

Authoritative statements: e.g. log records



- The basic CoSi protocol uses **Schnorr signatures**, that rely on a group G of prime order
 - *Discrete logarithmic problem is believed to be hard*

CoSi based on Schnorr Multisignature

- **Key Generation:**
 - Let G be a group of prime order r . Let g be a generator of G .
 - Select a random integer x in the interval $[0, r - 1]$. x is the private key and g^x is the public key.
 - N signers with individual private keys x_1, x_2, \dots, x_N , and the corresponding public keys $g^{x_1}, g^{x_2}, \dots, g^{x_N}$

CoSi based on Schnorr Multisignature

- **Signing:**

- Each signer picks up the random secret v_i , generates $V_i = g^{v_i}$
- The leader collects all such V_i , aggregates them $V = \prod V_i$, and uses a hash function to compute a collective challenge $c = H(V || S)$.

This challenge is forwarded to all the signers.

x_i is the private key. We sign using the private key.

- The signers send the response $r_i = v_i - cx_i$. The leader computes the aggregated as $r = \sum r_i$. The signature is (c, r) .

CoSi based on Schnorr Multisignature

- **Verification:**


- The verification key is $y = \prod g^{x_i}$ Gxi is the public key. We verify using public key.
- The signature is (c, r) , where $c = H(V||S)$ and $r = \sum r_i$
- Let $V_v = g^r y^c$
- Let $r_v = H(V_v||S)$
- If $r_v = r$, then the signature is verified

CoSi based on Schnorr Multisignature

- **Proof:**

- The verification key is $y = \prod g^{x_i}$
- The signature is (c, r) , where $c = H(V||S)$ and $r = \sum r_i$
- $V_v = g^r y^c = g^{\sum (v_i - cx_i)} \prod g^{cx_i} = g^{\sum (v_i - cx_i)} g^{\sum cx_i} = g^{\sum v_i} = \prod g^{v_i} = \prod V_i = V$
- So, $r_v = H(V_v||S) = H(V||S) = r$

CoSi Protocol



- One CoSi round to implement PBFT's pre-prepare and prepare phases
- Second CoSi round to implement PBFT's commit phase


Problems with Bitcoin

- There is **no verifiable commitment** of the system that a block would exist
 - Probability of successful fork attack decreases as the size of the Blockchain increases




Problems with Bitcoin

- There is **no verifiable commitment** of the system that a block would exist



Problems with Bitcoin-NG


- A faulty key block is verified only after end of the round
 - A faulty miner can introduce a number of correct microblocks following a faulty microblock in the system - certainly a overhead for the application - **a fork alleviates the problem further**



Solve this problem by a set of **PBFT verifier** - who will verify a block and then only the block is added in the Blockchain

PBFTCoin - A Strawman Design

- **Assumption:** $3f+1$ fixed "trustees" are there, who will run the PBFT to withstand f failures
 - Avoid the probabilistic strong consistency - introduces low latency in the system
 - No forks in the system - Blocks are added only after verification from the trustees




Problems of PBFT

- PBFT requires a **static consensus group** (because of message passing)
- **Scalability** (in terms of nodes) is a problem for PBFT
 - $O(n^2)$ communication complexity
 - $O(n)$ verification complexity
 - Absence of third-party verifiable proofs (PBFT uses MAC - need to share the keys among the miners)
- **Sybil attack** - create multiple pseudonymous identities to subvert the **$3f+1$** requirements of PBFT

Open the Consensus Group

- Use PoW based system to give a *proof of membership* of a miner as a part of the trustees
- Maintains a “balance of power” within the BFT consensus group
 - Use a fixed-size sliding window
 - Each time a miner finds a new block, it receives a *consensus group share*
 - The share proves the miner’s membership in the trustee group

Open the Consensus Group




Improve Efficiency

- **Improve $O(n)$ communication complexity**
 - Use tree based multicast protocol - share information with $O(\log n)$
- **Improve $O(n)$ complexity for verification**
 - Use Schnorr multisignatures or BLS for verification
 - Verification can be done in $O(1)$ through signature aggregation
- Multisignatures + Communication trees - **CoSi**

Further Improvement

- Inherit Bitcoin NG's idea of separating out **transaction verification** and **leader election**



ByzCoin Performance



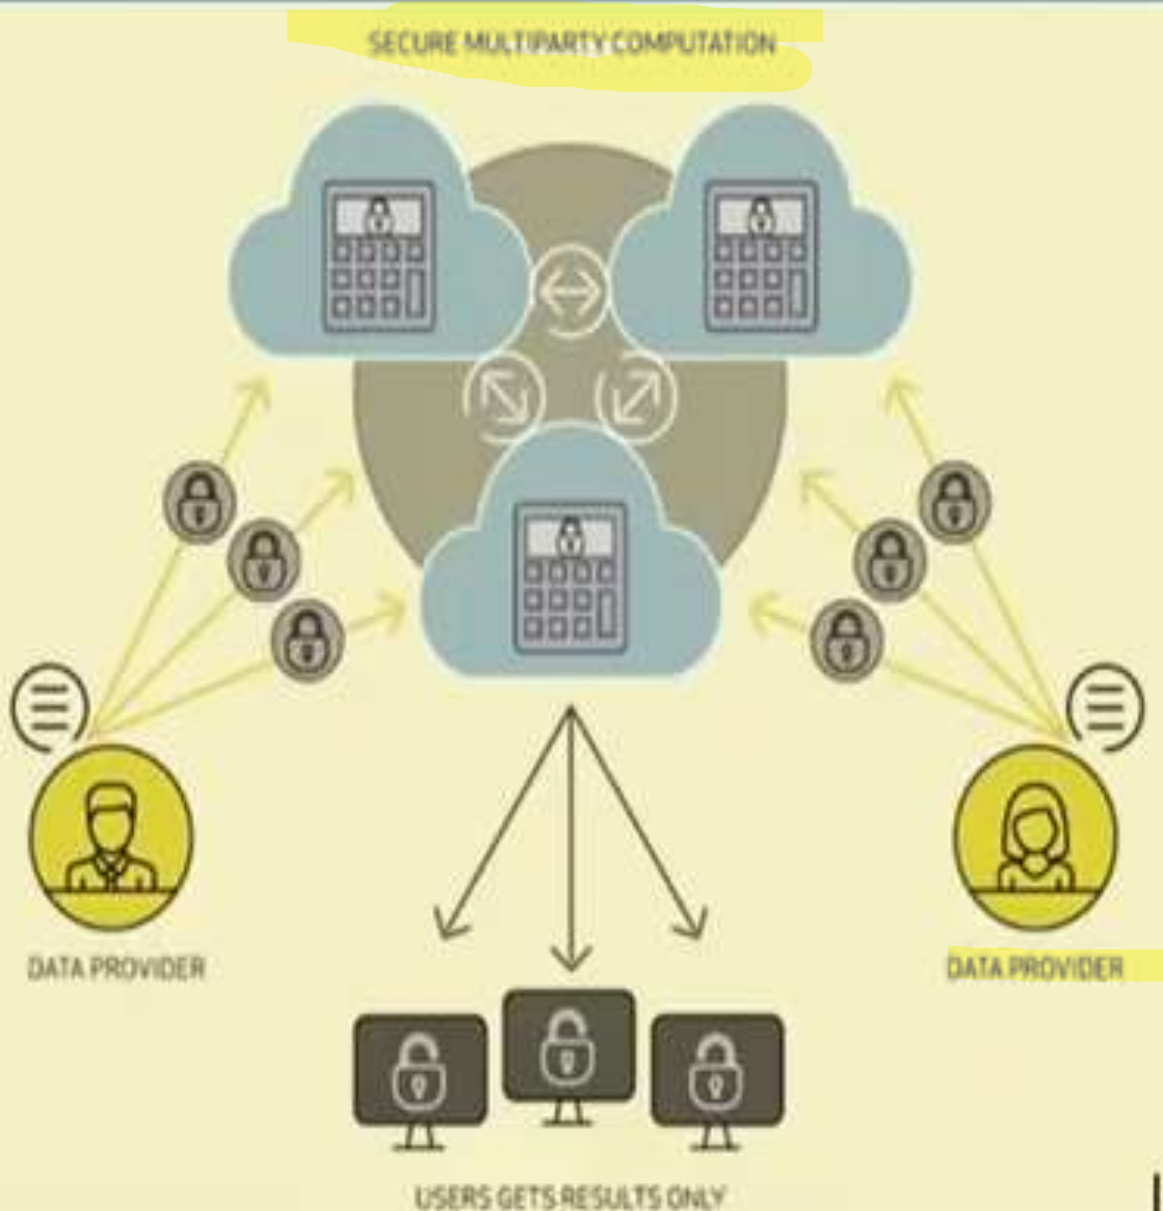


Image source: <https://thehub.dk/jobs/company/partisia>



Secured Multiparty Computation over Blockchain

Multiparty Computation (MPC)



- Information/data is distributed among multiple authorities with different data share or data distribution policies
- Users want to run a computation - however, the computation involves access to data from multiple sources

User doesn't want to reveal the private data, but wants to result of certain computation on the data

Dining Cryptographer Problem

- Three cryptographers are sitting down to dinner at their favorite restaurant
- Any of the cryptographer can pay the bill, or the bill can be directly paid by the National Security Council (NSC)



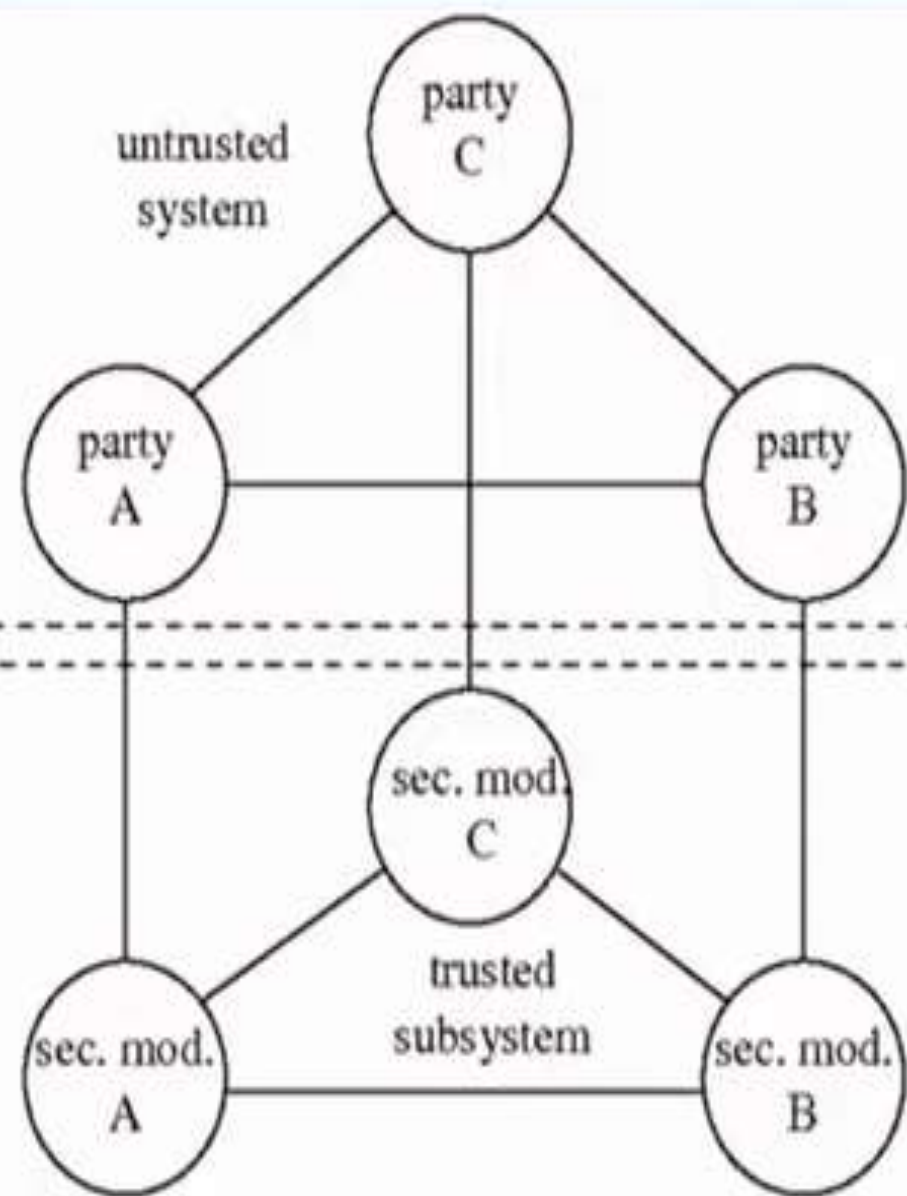
Dining Cryptographer Problem

- The three cryptographers respect each other's right to make an anonymous payment
 - But they wonder if NSA is paying
- This payment protocol can be designed using secured



Formal Definition

- There are n players p_1, p_2, \dots, p_n
- They wish to evaluate a function $f(x_1, x_2, \dots, x_n)$
- x_i is a secret value provided by p_i
- **Goal:**
 - Preserve the privacy of the player's input
 - Guarantee the correctness of the computation



Decentralized Solution

- The problem is trivial if we assume the pretense of a trusted third party - **but we do not want to have one**
- Two types of faulty behaviors in a decentralized system
 - Players may try to learn additional information (private computation)
 - Faulty players may try to disrupt the computation (secure computation)

Yao's Millionaire Problem

- Two millionaires wish to find out who is wealthier
- They do not want to reveal any other information



Preconditions

- We know the range of the inputs: $(0, N)$
- A: Public key **e**, Private key **d**
- B: Can access **e**, not **d**
- $D_d(E_e(X)) = X$
- $D_d(E_e(X) + Y) = \text{some random looking thing if you do not know } d$

Protocol Step 1

- A has i and B has j
- B generates a random x of m bits
- $C = E_e(X)$
- $u = C - (j - 1)$
- Send u to A

Protocol Step 2

- A Computes: *for* ($t = 1$ to N) $y_m = D_d(u+t)$
- A takes a prime p of size \sqrt{m} and computes
 - $z_i = y_i \bmod p$ for $i = 1$ to N
- p is chosen such that $|z_m - z_n| \geq 2$ for any m, n in $[1$ to $N]$

Protocol Step 3

- A sends B the following list
 - $p, z_1, z_2, \dots, z_i, (z_{i+1}+1), (z_{i+2}+1), \dots, (z_N+1)$
- B compares the j^{th} entry of this list excluding prime p with $(x \bmod p)$
- If $(x \bmod p) = j^{\text{th}}$ entry of the list, then $i \geq j$

Problems with MPC

- MPC has poor scaling properties
 - Performance in the malicious setting is worse than the semi-honest case
- Depends on the assumption that majority of the parties are always honest and share the correct information
 - How will you ensure that the parties have shared the correct information?
 - The parties can deny the sharing of information as well

Fair MPC

- Either all parties receive the protocol output or no party does
 - Extremely important for applications like auctions or contract signing
- Example:
 - Alice participates in an auction
 - She **learns first** that she did not win the auction
 - She aborts and claims a network failure - tries again with a new bid

Fair MPC

- [Cleve '86] Fair MPC is impossible to realize for general functions when a majority of the parties are dishonest.
 - Also holds when the parties have access to a trusted setup, such as a common reference string

Richard Cleve. Limits on the security of coin flips when half the processors are faulty (extended abstract). In STOC, pages 364–369, 1986

Solve Fair MPC - Use a Public Bulletin Board

- Parties have access to a *public ledger*
 - Allows anyone to publish arbitrary strings - used for MPC protocol
 - The strings contain proof about who has published the string - anyone can verify
- Run an unfair MPC protocol to compute an encryption of the function output - design a fair decryption protocol using the public ledger - either everyone can decrypt or no one can